

# Rossmann Store Sales Prediction Problem

*Shan Chen*

*Nov 27, 2015*

## Introduction

This is a pretty interesting problem cause I always wonder what can be the relatively vital effects to a store daily sales? My first intuitive on this I think might be weekends to be #1 and then possibly distance to competitors and then promotions (as far as I'm considered I was always "courageous" to buy those "buy one, get one" stuff), well, things would never be that easy.. So here we go, let's focus on the data and example to have an overall view on this prediction problem.

I'll basically cover the following procedures and those are kinda my thought process.

### 1. Feature Engineering and EDA

- 1.1 Data Preparation
- 1.2 Customers Vs. Sales
- 1.3 Open, StateHoliday Vs. Sales
- 1.4 DayOfWeek Vs. Sales
- 1.5 Date Vs. Sales
- 1.6 Competition Vs. Sales
- 1.7 Promotion Vs. Sales

### 2. Modeling and Prediction

- 2.1 Random Forest
- 2.2 Linear Model
- 2.3 Prediction

## 1. Feature Engineering and EDA

### 1.1 Data Preparation

Loading packages and data set

```
library('party')
```

```
## Loading required package: grid
```

```
## Loading required package: mvtnorm
```

```
## Loading required package: modeltools
```

```
## Loading required package: stats4
```

```
## Loading required package: strucchange
```

```
## Loading required package: zoo
```

```
##  
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':  
##  
##   as.Date, as.Date.numeric
```

```
## Loading required package: sandwich
```

```
library('rpart')  
library('randomForest') # classification algorithm
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
library('ggplot2') # Visualization
```

```
##  
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:randomForest':  
##  
##   margin
```

```
library('sqldf')
```

```
## Loading required package: gsubfn
```

```
## Loading required package: proto
```

```
## Loading required package: RSQLite
```

```
## Loading required package: DBI
```

```
library('dplyr') # Data manipulation
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:randomForest':  
##  
##      combine
```

```
## The following objects are masked from 'package:stats':  
##  
##      filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##      intersect, setdiff, setequal, union
```

```
library('lubridate') # Month variable extraction
```

```
##  
## Attaching package: 'lubridate'
```

```
## The following object is masked from 'package:base':  
##  
##      date
```

```
library('zoo')  
library('mice') # imputation
```

```
## Loading required package: Rcpp
```

```
## mice 2.25 2015-11-09
```

```
train<-read.csv("C:\\Users\\Shan\\Desktop\\Rossmann Store Sales Prediction\\train.csv")  
store<-read.csv("C:\\Users\\Shan\\Desktop\\Rossmann Store Sales Prediction\\store.csv")  
test<-read.csv("C:\\Users\\Shan\\Desktop\\Rossmann Store Sales Prediction\\test.csv")
```

train(test) and store are two related data sets with the same store id column, here let's use `sqldf()` to horizontally join them to get a complete set.

```
#first vertically bind train and test sets to get a complete set  
complete<- bind_rows(train,test)
```

```
## Warning in rbind_all(x, .id): Unequal factor levels: coercing to character
```

```
## Warning in rbind_all(x, .id): Unequal factor levels: coercing to character
```

```
complete<- sqldf("select complete.*,store.* from complete  
inner join store on complete.store=store.store")
```

```
## Loading required package: tcltk
```

```
dim(train)
```

```
## [1] 1017209      9
```

```
dim(test)
```

```
## [1] 41088      8
```

```
dim(store)
```

```
## [1] 1115     10
```

```
complete<-complete[,-11]  
train<-complete[1:1017209,]  
test<-complete[1017210:1058279,]  
dim(complete)
```

```
## [1] 1058297     19
```

```
head(complete)
```

##	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	StateHoliday
## 1	1	5	7/31/2015	5263	555	1	1	0
## 2	2	5	7/31/2015	6064	625	1	1	0
## 3	3	5	7/31/2015	8314	821	1	1	0
## 4	4	5	7/31/2015	13995	1498	1	1	0
## 5	5	5	7/31/2015	4822	559	1	1	0
## 6	6	5	7/31/2015	5651	589	1	1	0

##	SchoolHoliday	Id	StoreType	Assortment	CompetitionDistance
## 1	1	NA	c	a	1270
## 2	1	NA	a	a	570
## 3	1	NA	a	a	14130
## 4	1	NA	c	c	620
## 5	1	NA	a	a	29910
## 6	1	NA	a	a	310

##	CompetitionOpenSinceMonth	CompetitionOpenSinceYear	Promo2
## 1	9	2008	0
## 2	11	2007	1
## 3	12	2006	1
## 4	9	2009	0
## 5	4	2015	0
## 6	12	2013	0

##	Promo2SinceWeek	Promo2SinceYear	PromoInterval
## 1	NA	NA	
## 2	13	2010	Jan, Apr, Jul, Oct
## 3	14	2011	Jan, Apr, Jul, Oct
## 4	NA	NA	
## 5	NA	NA	
## 6	NA	NA	

```
str(complete)
```


```
## 'data.frame': 1058297 obs. of 19 variables:
## $ Store : int 1 2 3 4 5 6 7 8 9 10 ...
## $ DayOfWeek : int 5 5 5 5 5 5 5 5 5 5 ...
## $ Date : chr "7/31/2015" "7/31/2015" "7/31/2015" "7/31/2015"
...
## $ Sales : int 5263 6064 8314 13995 4822 5651 15344 8492 8565 71
85 ...
## $ Customers : int 555 625 821 1498 559 589 1414 833 687 681 ...
## $ Open : int 1 1 1 1 1 1 1 1 1 1 ...
## $ Promo : int 1 1 1 1 1 1 1 1 1 1 ...
## $ StateHoliday : chr "0" "0" "0" "0" ...
## $ SchoolHoliday : int 1 1 1 1 1 1 1 1 1 1 ...
## $ Id : int NA NA NA NA NA NA NA NA NA NA ...
## $ StoreType : Factor w/ 4 levels "a","b","c","d": 3 1 1 3 1 1 1 1 1 1
1 ...
## $ Assortment : Factor w/ 3 levels "a","b","c": 1 1 1 3 1 1 3 1 3 1
...
## $ CompetitionDistance : int 1270 570 14130 620 29910 310 24000 7520 2030 3160
...
## $ CompetitionOpenSinceMonth: int 9 11 12 9 4 12 4 10 8 9 ...
## $ CompetitionOpenSinceYear : int 2008 2007 2006 2009 2015 2013 2013 2014 2000 2009
...
## $ Promo2 : int 0 1 1 0 0 0 0 0 0 0 ...
## $ Promo2SinceWeek : int NA 13 14 NA NA NA NA NA NA NA ...
## $ Promo2SinceYear : int NA 2010 2011 NA NA NA NA NA NA NA ...
## $ PromoInterval : Factor w/ 4 levels "", "Feb,May,Aug,Nov",...: 1 3 3 1 1
1 1 1 1 1 ...
```

Dimensionality for complete passed through our checking, next let's consider those predictors.

## 1.2 Customers Vs. Sales

```
ggplot(train, aes(x=Customers, y=Sales)) + geom_point()+ggtitle("Scatter plot of Sales
and Customers")
```

Scatter plot of Sales and Customers



No doubt that Customers is one perfect predictor for sales, however since we can not know the future Customers, I'm not using Customers in the prediction.

## 1.3 Open, StateHoliday Vs. Sales

```
a=table(complete$Open, complete$StateHoliday)
addmargins(a)
```

```
##
##           0         a         b         c         Sum
##  0    148507    19720    6545    4029    178801
##  1    878549     720     145     71    879485
## Sum 1027056    20440    6690    4100 1058286
```

There are 178801 out of 1058286 obs are closed so have no sales, it may due to holiday or for other private reasons. Hereby I will not include those obs with closed days. Though mostly all of the stores would be closed on the holiday (what a life!), there are still some rare stores keep opening, let's do a boxplot of Sales on StateHoliday.


```
#Delete rows of Sales="0"
train<-sqldf("select * from train where Sales!=0 ")
test<-sqldf("select * from test where Open=1 ")
complete<- bind_rows(train,test)
dim(train)
```

```
## [1] 844338    19
```

```
dim(test)
```

```
## [1] 35075    19
```

```
#Subset complete to new train and test sets
train<-complete[1:844338,]
test<-complete[844339:879413,]
#Factorize StateHoliday Variable
complete$SchoolHoliday <- as.factor(complete$SchoolHoliday)
complete$StateHoliday<-as.factor(complete$StateHoliday)
#Boxplots
Sales_vs_StateHoliday <- ggplot(train, aes(x=StateHoliday, y=Sales)) + geom_boxplot()
Sales_vs_StateHoliday
```




It seems that Sales are higher on holidays.

## 1.4 DayOfWeek Vs. Sales



```
#Factorize DayOfWeek Variable
complete$DayOfWeek<-as.factor(complete$DayOfWeek)
#Boxplots
Sales_vs_DayOfWeek <- ggplot(complete[1:844338,], aes(x=DayOfWeek, y=Sales)) + geom_box
plot()
Sales_vs_DayOfWeek
```



```
#total_Counts for DayOfWeek
total_Counts<-sqldf("select DayOfWeek, count(*) from train group by DayOfWeek order by
DayOfWeek")
total_Counts
```

```
## DayOfWeek count(*)
## 1          1  137557
## 2          2  143955
## 3          3  141922
## 4          4  134626
## 5          5  138633
## 6          6  144052
## 7          7   3593
```

From `total_Counts` lists, the number of obs for Sunday is much lower than the other week days possibly due to Germany tend to close all the business on Sunday, but the sales tend to be higher, same with Monday which is quite interesting.


## 1.5 Date Vs. Sales

For variable of Date, I would like to extract the month information out (since DayOfWeek is preferred than the day information from Date and year is too overall for our daily sales prediction).

```
#as.Date Date variable
complete$Date=as.Date(complete$Date,format = "%m/%d/%Y")
complete=complete %>%
  filter%>%
  mutate(Month = lubridate::month(Date))
```

Boxplots to show effects of new added variable Month on Sales

```
#Factorize Month Variable
complete$Month<-as.factor(complete$Month)
#Boxplots
Sales_vs_Month <- ggplot(complete[1:844338,], aes(x=Month, y=Sales)) + geom_boxplot()
Sales_vs_Month
```



It seems Dec has the highest Sales, it's obvious that people all love to shopping at the end of the year.

## 1.6 Competition Vs. Sales


Competition relationship at my first sight should be one important factor to our prediction, stores with rare(too far away) competitors may have a relatively higher sales than the others, but the sales possibly would not be that high due to also lack of population of residence around...It could be a

complicated point to analyze thoughly, maybe we should combine the number of customers and competiton information to show a kinda significant factor variable.

First let's calculate years since opening of the competitor store.

```
complete$CompetitionOpenYears <- as.yearmon("2015-07-31") - as.yearmon(paste(complete$CompetitionOpenSinceYear,complete$CompetitionOpenSinceMonth, sep = "-"))  
#Histogram  
hist(complete$CompetitionOpenYears,breaks=100,main = "Years since opening of Competitor")
```

### Years since opening of Competitor



```
#many missing values  
summary(complete$CompetitionOpenYears)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's  
## -0.08    2.42    5.42    6.29   9.17   115.50  281598
```

```
#Impute N/A's with mean  
complete[is.na(complete$CompetitionOpenYears), c("CompetitionOpenYears")] = mean(complete$CompetitionOpenYears,na.rm = TRUE)
```

Then go to Competition\$Distance part

```
summary(complete$CompetitionDistance)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	20	710	2330	5446	6880	75860	2251

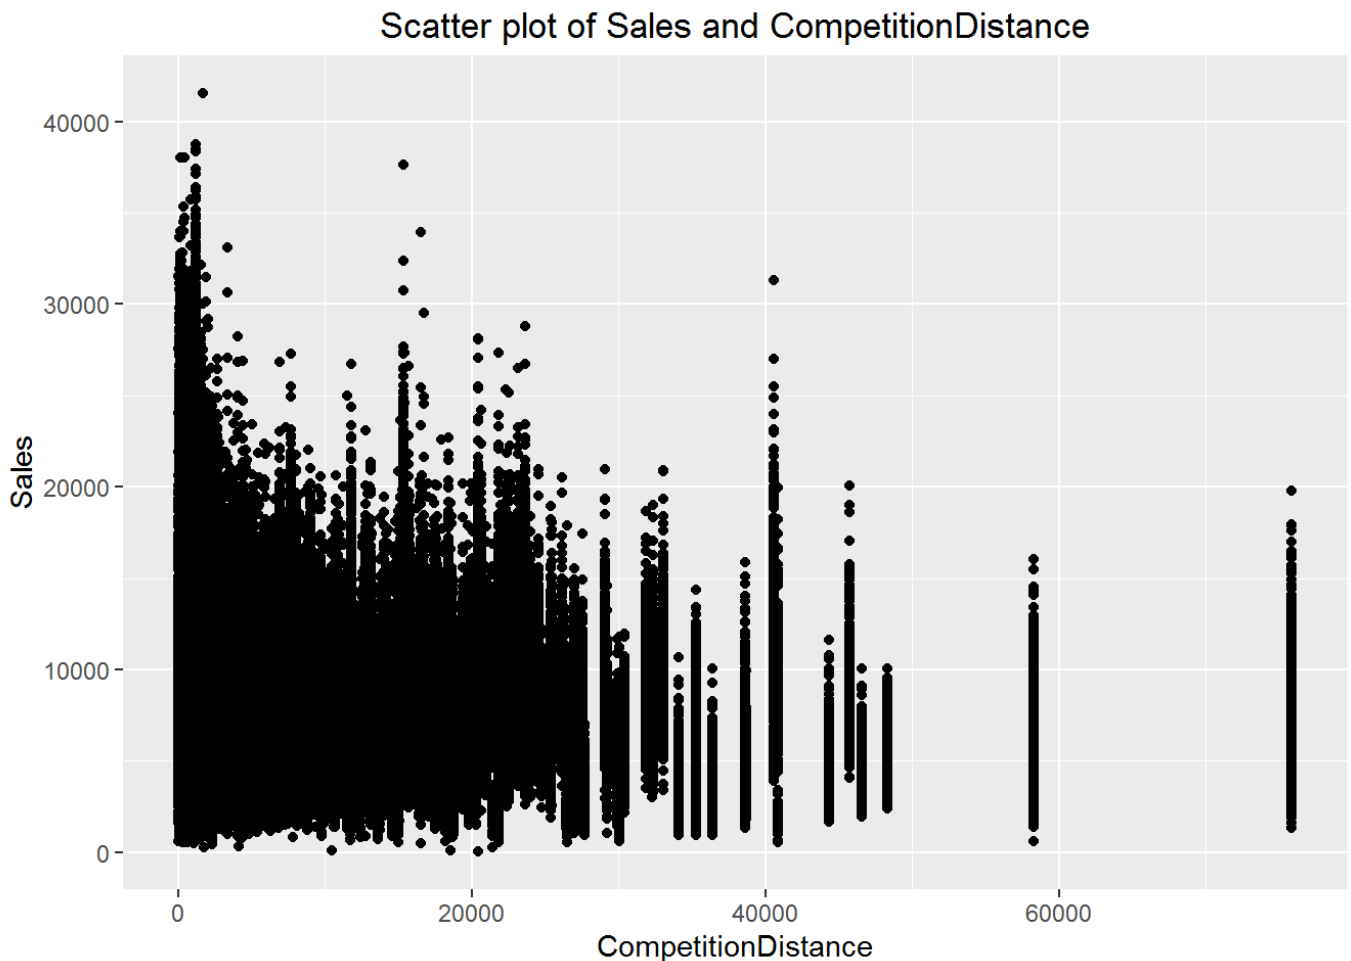
*#Three stores have missing CompetitionDistance value*

Impute the N/A CompetitionDistance observations with the mean of CompetitionDistance observations.

```
complete[is.na(complete$CompetitionDistance), c("CompetitionDistance")] = mean(complete$CompetitionDistance, na.rm = TRUE)
```

Scatterplot:

```
ggplot(complete[1:844338,], aes(x=CompetitionDistance, y=Sales)) + geom_point()+ggtitle("Scatter plot of Sales and CompetitionDistance")
```



It's hard to describe the pattern of the two variables, but we can still tell that **those highest Daily Sales did happen on the stores with the nearest competitor**, which is not what I thought about the competition effects(nearer the competitor is, lower the Sales would be), so such situation may happen due to the corresponding crowded population around.

one way I'll use to try to split the effects is to split the population crowd to four levels through **discreting Customers variable**.

```
#build a new data set with store id and mean number of customers visited that store before
```

```
mean_customers_store <- sqldf("select Store, avg(Customers) as MeanCustomers from train  
group by Store order by Store")  
head(mean_customers_store)
```

```
##   Store MeanCustomers  
## 1     1      564.0499  
## 2     2      583.9987  
## 3     3      750.0770  
## 4     4     1321.7526  
## 5     5      537.3402  
## 6     6      635.2346
```

```
#get quatiles of MeanCustomers as the splition criterion  
summary(mean_customers_store$MeanCustomers)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
##  240.2   541.5   678.7   754.6   866.2  3403.0
```

```
#Creat a new Customers dependent variable called PopulationLevel
```


```
mean_customers_store$PopulationLevel <- 1  
mean_customers_store$PopulationLevel[mean_customers_store$MeanCustomers >541.5 & mean_c  
ustomers_store$MeanCustomers <=678.7 ] <- 2  
mean_customers_store$PopulationLevel[mean_customers_store$MeanCustomers >678.7 & mean_c  
ustomers_store$MeanCustomers <=866.2 ] <- 3  
mean_customers_store$PopulationLevel[mean_customers_store$MeanCustomers >866.2 ] <- 4
```

```
#Join PopulationLevel into our complete data
```

```
complete<- sqldf("select complete.*,b.PopulationLevel from complete  
                inner join mean_customers_store b on complete.Store=b.Store")
```

```
#Density plot of Sales on PopulationLevel
```


```
complete$PopulationLevel <- as.factor(complete$PopulationLevel)  
ggplot(complete[1:844338,], aes(x=Sales, color=PopulationLevel)) + geom_density() + ggt  
itle("Density plot of Sales on PopulationLevel")
```




The multiple density plots show Sales is more right skewed on the higher population level. Here, we again take in the important predictor Customers' information, it's kinda addressed the issue of population size nearby.

## 1.7 Promotion Vs. Sales

```
#Factorize Promo Variable
complete$Promo<-as.factor(complete$Promo)
complete$Promo2<-as.factor(complete$Promo2)
#Boxplots
ggplot(complete[1:844338,], aes(x=Promo, y=Sales)) + geom_boxplot()
```



```
ggplot(complete[1:844338,], aes(x=Promo2, y=Sales)) + geom_boxplot()
```




For Promotion running on that day, the sales would increase a bit, but not the same for the consecutive promotion activity. Lets guess that it's possible only poor sales store need a knida long-term promotion to be more attractive to customers.. it makes sense!

Next, let's consider the duration of Promo2.

```
#Creat new variable complete$Promo2OpenDays
complete$Promo2OpenDays <- as.numeric(as.POSIXct("2015-07-31", format = "%Y-%m-%d") - a
s.POSIXct(paste(complete$Promo2SinceYear,complete$Promo2SinceWeek, 1, sep = "-"),format
= "%Y-%U-%u"))
#histogram
hist(complete$Promo2OpenDays, breaks=100, main = "Days since opening of promo2")
```



## Days since opening of promo2




```
#impute NA's with 0 days
complete[is.na(complete$Promo2OpenDays), c("Promo2OpenDays")] = 0
summary(complete$Promo2OpenDays)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##         0         0      53    620   1215   2188
```

Finally for promo2interval, there are four levels None , Feb,May,Aug,Nov , Jan,Apr,Jul,Oct and Mar,Jun,Sept,Dec . From the boxplots below, it seems sales differ cross different promo-renew-interval.

```
ggplot(complete[1:844338,], aes(x=PromoInterval, y=Sales)) + geom_boxplot()
```



## 2. Modeling and Prediction

```
#check data type and missing value
complete<-sqldf("select Id,Sales,DayOfWeek,Month,StateHoliday,SchoolHoliday,StoreType,Assortment,CompetitionDistance,CompetitionOpenYears,PopulationLevel,Promo,Promo2,Promo2OpenDays,PromoInterval from complete")
str(complete)
```

```
## 'data.frame': 879413 obs. of 15 variables:
## $ Id : int NA NA NA NA NA NA NA NA NA ...
## $ Sales : int 5263 6064 8314 13995 4822 5651 15344 8492 8565 7185
...
## $ DayOfWeek : Factor w/ 7 levels "1","2","3","4",...: 5 5 5 5 5 5 5 5 5 5
...
## $ Month : Factor w/ 12 levels "1","2","3","4",...: 7 7 7 7 7 7 7 7 7 7
...
## $ StateHoliday : Factor w/ 4 levels "0","a","b","c": 1 1 1 1 1 1 1 1 1 1 ...
## $ SchoolHoliday : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 2 2 2 ...
## $ StoreType : Factor w/ 4 levels "a","b","c","d": 3 1 1 3 1 1 1 1 1 1 ...
## $ Assortment : Factor w/ 3 levels "a","b","c": 1 1 1 3 1 1 3 1 3 1 ...
## $ CompetitionDistance : num 1270 570 14130 620 29910 ...
## $ CompetitionOpenYears: num 6.83 7.67 8.58 5.83 0.25 ...
## $ PopulationLevel : Factor w/ 4 levels "1","2","3","4": 2 2 3 4 1 2 4 2 2 2 ...
## $ Promo : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 2 2 2 ...
## $ Promo2 : Factor w/ 2 levels "0","1": 1 2 2 1 1 1 1 1 1 1 ...
## $ Promo2OpenDays : num 0 1950 1579 0 0 ...
## $ PromoInterval : Factor w/ 4 levels "", "Feb,May,Aug,Nov",...: 1 3 3 1 1 1 1 1 1 1
1 1 ...
```

```
#Give "No Interval" to missing value in PromoInterval
sqldf() # start a sequence of SQL statements
```

```
## <SQLiteConnection>
```

```
fn$sqldf("update complete set PromoInterval = 'No Interval' where PromoInterval = '' ")
```

```
## NULL
```

```
complete<- sqldf("select * from main.complete")
sqldf() # SQL statements finished
```

```
## NULL
```

```
#Factorization all cha columns
complete<- as.data.frame(unclass(complete))
summary(complete$Sales)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##      46    4859    6369    6956    8360    41550   35075
```

```
summary(complete$CompetitionDistance)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      20     710    2330    5446    6880    75860
```

```
summary(complete$CompetitionOpenYears)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## -0.08333  3.66700   6.28600   6.28600   7.25000  115.50000
```

```
complete$CompetitionOpenYears[complete$CompetitionOpenYears<0]<-0
summary(complete$Promo20OpenDays)
```

```
##      Min. 1st Qu.  Median     Mean 3rd Qu.    Max.
##       0       0      53      620   1215   2188
```

## 2.1 Random Forest

In this session, we are ready to predict Sales for Rossmann stores based on variables that we carefully treated and extracted. First let's use `randomForest` classification algorithm to deal with such a bunch of categorical variables.

```
#split complete to train and test data set
train<-complete[1:844338,-1]
test<-complete[844339:879413,]

# Using a random sample from the train set due to the PC memory issue
#I first run a sample of 10,000, and set the ntree option to 100
sample<-train[sample(nrow(train), 10000), ]

#Build rf model
rf.model1 <- randomForest(Sales~DayOfWeek+Month+StateHoliday+SchoolHoliday+StoreType+Assortment+CompetitionDistance+CompetitionOpenYears+PopulationLevel+Promo+Promo2+Promo20OpenDays+PromoInterval,data=sample,importance=TRUE, proximity=TRUE,ntree=100)
rf.model1
```

```
##
## Call:
## randomForest(formula = Sales ~ DayOfWeek + Month + StateHoliday + SchoolHoliday + StoreType + Assortment + CompetitionDistance + CompetitionOpenYears + PopulationLevel + Promo + Promo2 + Promo20OpenDays + PromoInterval, data = sample, importance = TRUE, proximity = TRUE, ntree = 100)
##              Type of random forest: regression
##              Number of trees: 100
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 3002073
##              % Var explained: 67.9
```

*#Then I tried sample size of 15,000*

```
sample<-train[sample(nrow(train), 15000), ]
```

```
rf.model2 <- randomForest(Sales~DayOfWeek+Month+StateHoliday+SchoolHoliday+StoreType+Assortment+CompetitionDistance+CompetitionOpenYears+PopulationLevel+Promo+Promo2+Promo2OpenDays+PromoInterval,data=sample,importance=TRUE, proximity=TRUE,ntree=100)
rf.model2
```

```
##
```

```
## Call:
```

```
## randomForest(formula = Sales ~ DayOfWeek + Month + StateHoliday + SchoolHoliday + StoreType + Assortment + CompetitionDistance + CompetitionOpenYears + PopulationLevel + Promo + Promo2 + Promo2OpenDays + PromoInterval, data = sample, importance = TRUE, proximity = TRUE, ntree = 100)
```

```
##           Type of random forest: regression
```

```
##           Number of trees: 100
```

```
## No. of variables tried at each split: 4
```

```
##
```

```
##           Mean of squared residuals: 3047404
```

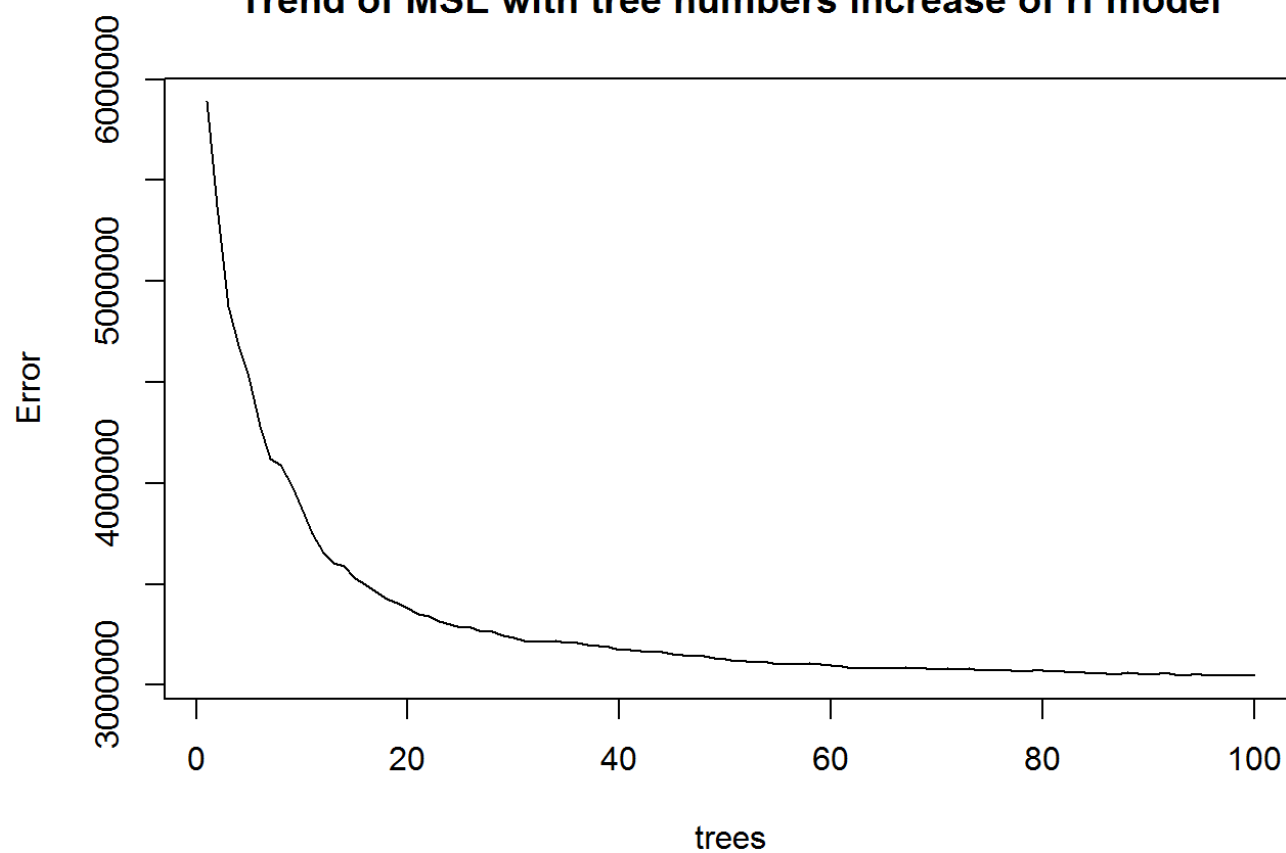
```
##           % Var explained: 68.81
```

*#Won't work anymore for the size of 20,000, I need to buy a bigger computer :(*

*# Show model error, since my response variable Sales is one continuous variable, the randomForest is running a regression type and the error rate plot would only show one black solid line representing the OOB MSE*

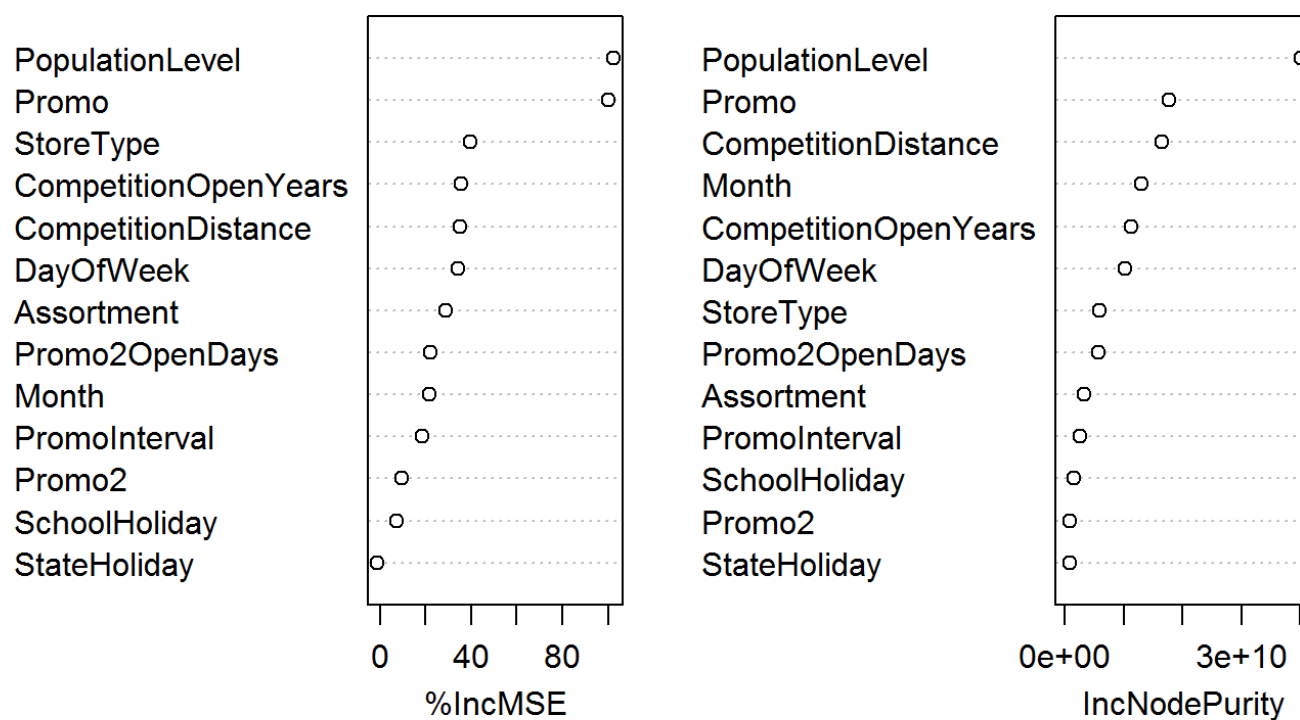
```
plot(rf.model2,main="Trend of MSE with tree numbers increase of rf model")
```

**Trend of MSE with tree numbers increase of rf model**



```
#Show the relative variable importance  
varImpPlot(rf.model2, main="Relative Variable Importance")
```

## Relative Variable Importance

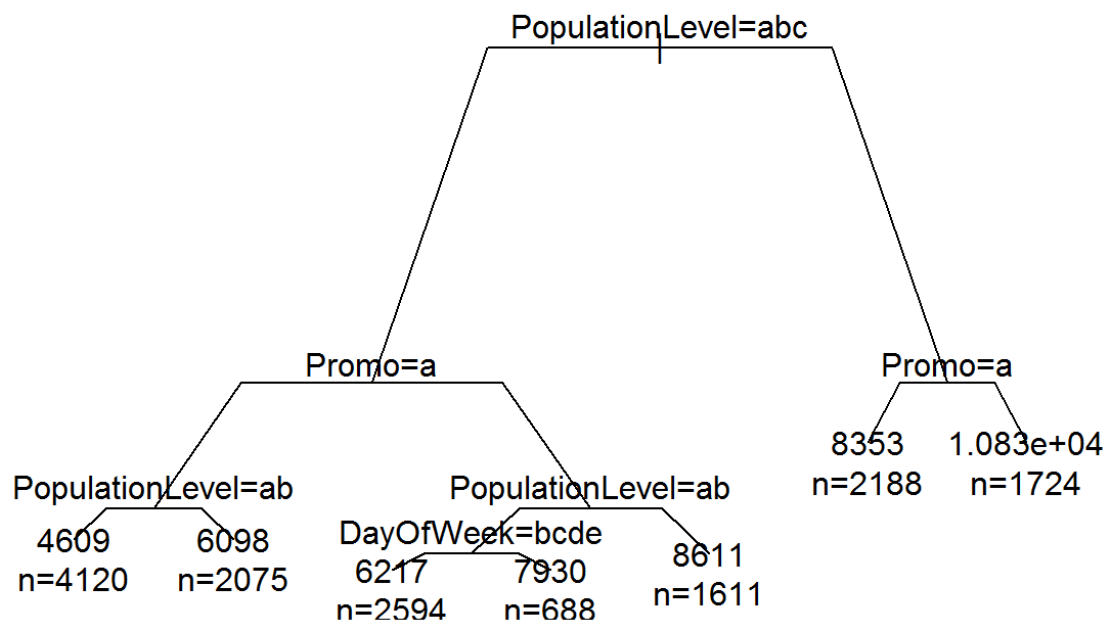


*#Build rpart the model*

```
rpart.model <- rpart(Sales~DayOfWeek+Month+StateHoliday+SchoolHoliday+StoreType+Assortment+CompetitionDistance+CompetitionOpenYears+PopulationLevel+Promo+Promo2+Promo2OpenDays+PromoInterval,data=sample)
```

```
plot(rpart.model, branch=0.6, margin=0.05)
```

```
text(rpart.model, use.n=TRUE)
```



```
summary(rpart.model)
```



```

## Call:
## rpart(formula = Sales ~ DayOfWeek + Month + StateHoliday + SchoolHoliday +
##       StoreType + Assortment + CompetitionDistance + CompetitionOpenYears +
##       PopulationLevel + Promo + Promo2 + Promo2OpenDays + PromoInterval,
##       data = sample)
##       n= 15000
##
##           CP nsplit rel error      xerror      xstd
## 1 0.22735527      0 1.0000000 1.0000310 0.02238110
## 2 0.08532602      1 0.7726447 0.7727562 0.01719501
## 3 0.04045673      2 0.6873187 0.6874735 0.01696451
## 4 0.03053721      3 0.6468620 0.6472056 0.01668164
## 5 0.02088304      4 0.6163248 0.6167061 0.01656095
## 6 0.01088748      5 0.5954417 0.5958872 0.01639407
## 7 0.01000000      6 0.5845543 0.5850116 0.01632627
##
## Variable importance
##       PopulationLevel          Promo          DayOfWeek
##                58                26                6
## CompetitionDistance      StoreType      Assortment
##                4                3                2
##
## Node number 1: 15000 observations,      complexity param=0.2273553
## mean=6936.563, MSE=9771023
## left son=2 (11088 obs) right son=3 (3912 obs)
## Primary splits:
##       PopulationLevel splits as LLLR, improve=0.22735530, (0 missing)
##       Promo           splits as LR, improve=0.12504050, (0 missing)
##       DayOfWeek       splits as RLLLLLR, improve=0.02908482, (0 missing)
##       Promo2OpenDays < 598.9583 to the right, improve=0.01981068, (0 missing)
##       Month           splits as LLLRLLLLLLLL, improve=0.01944527, (0 missing)
## Surrogate splits:
##       CompetitionDistance < 275 to the right, agree=0.758, adj=0.073, (0 split)
##       StoreType           splits as LRLL, agree=0.758, adj=0.071, (0 split)
##       Assortment          splits as LRL, agree=0.749, adj=0.037, (0 split)
##       DayOfWeek           splits as LLLLLLR, agree=0.742, adj=0.011, (0 split)
##       StateHoliday        splits as LRRR, agree=0.740, adj=0.002, (0 split)
##
## Node number 2: 11088 observations,      complexity param=0.08532602
## mean=6051.253, MSE=5314217
## left son=4 (6195 obs) right son=5 (4893 obs)
## Primary splits:
##       Promo           splits as LR, improve=0.21223660, (0 missing)
##       PopulationLevel splits as LLR-, improve=0.12288250, (0 missing)
##       DayOfWeek       splits as RLLLLLL, improve=0.05019538, (0 missing)
##       StoreType       splits as L-LR, improve=0.02829672, (0 missing)
##       Month           splits as LLLRLLLLLLLL, improve=0.02396391, (0 missing)
## Surrogate splits:
##       DayOfWeek           splits as RRRRRLL, agree=0.614, adj=0.126, (0 split)
##       StateHoliday        splits as LR--, agree=0.559, adj=0.001, (0 split)
##       CompetitionOpenYears < 0.04166667 to the right, agree=0.559, adj=0.000, (0 split)

```

```

##
## Node number 3: 3912 observations,      complexity param=0.04045673
## mean=9445.846, MSE=1.38852e+07
## left son=6 (2188 obs) right son=7 (1724 obs)
## Primary splits:
##      Promo                splits as LR, improve=0.10916190, (0 missing)
##      DayOfWeek            splits as RLLLLLR, improve=0.02957379, (0 missing)
##      Month                splits as LLLRLLLLLLLL, improve=0.02639026, (0 missing)
##      Assortment           splits as LLR, improve=0.01510439, (0 missing)
##      CompetitionOpenYears < 2.208333 to the right, improve=0.01213442, (0 missin
g)
## Surrogate splits:
##      DayOfWeek            splits as RRRRRLL, agree=0.626, adj=0.151, (0 split)
##      CompetitionOpenYears < 22.83333 to the left, agree=0.561, adj=0.003, (0 spl
it)
##      CompetitionDistance < 58200 to the left, agree=0.560, adj=0.001, (0 spl
it)
##      StateHoliday         splits as LLRL, agree=0.560, adj=0.001, (0 split)
##
## Node number 4: 6195 observations,      complexity param=0.02088304
## mean=5107.417, MSE=3417209
## left son=8 (4120 obs) right son=9 (2075 obs)
## Primary splits:
##      PopulationLevel      splits as LLR-, improve=0.14458140, (0 missing)
##      StoreType            splits as L-LR, improve=0.05324117, (0 missing)
##      Assortment           splits as L-R, improve=0.04207632, (0 missing)
##      CompetitionDistance < 2080 to the left, improve=0.03820283, (0 missing)
##      Month                splits as LLLRLLLLLLLL, improve=0.03301390, (0 missing)
## Surrogate splits:
##      CompetitionDistance < 205 to the right, agree=0.693, adj=0.082, (0 spl
it)
##      CompetitionOpenYears < 16.04167 to the left, agree=0.669, adj=0.011, (0 spl
it)
##      DayOfWeek            splits as LLLLLLR, agree=0.666, adj=0.001, (0 split)
##      StateHoliday         splits as LR--, agree=0.665, adj=0.000, (0 split)
##
## Node number 5: 4893 observations,      complexity param=0.03053721
## mean=7246.239, MSE=5160145
## left son=10 (3282 obs) right son=11 (1611 obs)
## Primary splits:
##      PopulationLevel      splits as LLR-, improve=0.17726520, (0 missing)
##      DayOfWeek            splits as RLLLL--, improve=0.11007010, (0 missing)
##      Month                splits as LLLRLLLLLLLL, improve=0.03867596, (0 missing)
##      StoreType            splits as L-LR, improve=0.02173866, (0 missing)
##      CompetitionDistance < 3855 to the left, improve=0.02150701, (0 missing)
## Surrogate splits:
##      CompetitionDistance < 145 to the right, agree=0.692, adj=0.065, (0 spl
it)
##      CompetitionOpenYears < 16.04167 to the left, agree=0.674, adj=0.010, (0 spl
it)
##
## Node number 6: 2188 observations
## mean=8353.006, MSE=1.193559e+07
##

```

```

## Node number 7: 1724 observations
##   mean=10832.81, MSE=1.292011e+07
##
## Node number 8: 4120 observations
##   mean=4608.587, MSE=2520577
##
## Node number 9: 2075 observations
##   mean=6097.864, MSE=3722459
##
## Node number 10: 3282 observations,   complexity param=0.01088748
##   mean=6576.168, MSE=3803335
##   left son=20 (2594 obs) right son=21 (688 obs)
##   Primary splits:
##       DayOfWeek      splits as  RLLLL--, improve=0.12783670, (0 missing)
##       PopulationLevel splits as  LR--, improve=0.09620262, (0 missing)
##       StoreType      splits as  L-LR, improve=0.05847065, (0 missing)
##       Month          splits as  LLLRLLLLLLLL, improve=0.05188765, (0 missing)
##       Promo2OpenDays < 2117.979 to the left, improve=0.02733961, (0 missing)
##
## Node number 11: 1611 observations
##   mean=8611.338, MSE=5146088
##
## Node number 20: 2594 observations
##   mean=6217.064, MSE=2938934
##
## Node number 21: 688 observations
##   mean=7930.112, MSE=4743060

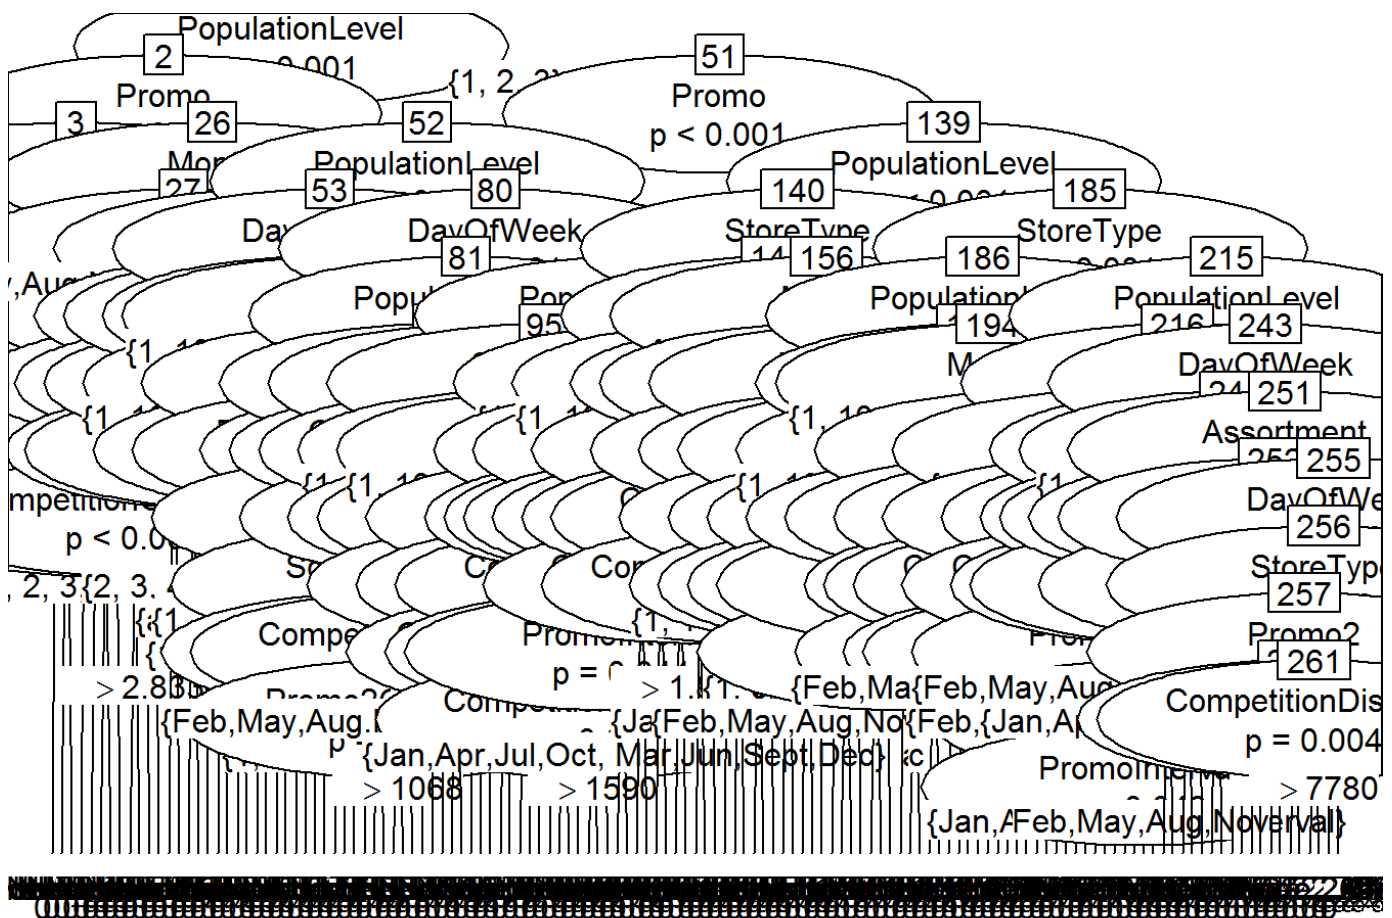
```

*#Finally I tried Conditional Inference Tree*

```

ctree.model <- ctree(Sales~DayOfWeek+Month+StateHoliday+SchoolHoliday+StoreType+Assortm
ent+CompetitionDistance+CompetitionOpenYears+PopulationLevel+Promo+Promo2+Promo2OpenDay
s+PromoInterval,data=sample)
plot(ctree.model)

```



See? the OOB error rate from rf though is not so pretty, around 0.3, as long as I increased the sample size from 10,000 to 15,000 and keep ntree the same, the error rate would decrease as well. So I would like to conclude that our model is still reasonable once there's a way to include that really large train data set(844,392 obs).

Let's have a look at the Relative Variable Importance table above. Whoa, NICE we made our PopulationLevel Variable(extracted information from our vital predictor Customers), and as we guess earlier that Promo and Competition gave some importance to the store Sales. Also some other new created variable such as CompetitionOpenYears, Month and Promo2OpenDays are all ranked well, from this case we can tell that feature engineering work is so worthing!

For regression tree model, I added another two algorithms: rpart and ctree for comparison

## 2.2 Linear Model

Just out of curiosity, I still wanna run a simple linear model including the same features, maybe just as a comparison of my rf model.

```
#Linear Model
lm.fit <- lm(Sales~.,data=train)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = Sales ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11478.2  -1254.2   -185.0    948.2   31619.6
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error  t value Pr(>|t|)
## (Intercept)    3.785e+03  1.200e+01  315.434 < 2e-16 ***
## DayOfWeek2     -1.077e+03  8.124e+00 -132.517 < 2e-16 ***
## DayOfWeek3     -1.404e+03  8.157e+00 -172.155 < 2e-16 ***
## DayOfWeek4     -1.371e+03  8.270e+00 -165.799 < 2e-16 ***
## DayOfWeek5     -1.013e+03  8.207e+00 -123.416 < 2e-16 ***
## DayOfWeek6     -9.672e+02  8.778e+00 -110.179 < 2e-16 ***
## DayOfWeek7     -9.850e+02  3.828e+01  -25.731 < 2e-16 ***
## Month10        -1.904e+01  1.194e+01   -1.595    0.111
## Month11         4.752e+02  1.202e+01   39.517 < 2e-16 ***
## Month12         1.965e+03  1.211e+01  162.279 < 2e-16 ***
## Month2          8.597e+01  1.056e+01    8.140 3.96e-16 ***
## Month3          3.019e+02  1.038e+01   29.086 < 2e-16 ***
## Month4          4.259e+02  1.055e+01   40.377 < 2e-16 ***
## Month5          5.249e+02  1.057e+01   49.636 < 2e-16 ***
## Month6          4.461e+02  1.049e+01   42.524 < 2e-16 ***
## Month7          1.792e+02  1.061e+01   16.894 < 2e-16 ***
## Month8         -9.885e+01  1.227e+01   -8.058 7.77e-16 ***
## Month9         -8.032e+01  1.194e+01   -6.728 1.73e-11 ***
## StateHolidaya   -3.458e+01  8.224e+01   -0.420    0.674
## StateHolidayb   -2.665e+02  1.794e+02   -1.485    0.138
## StateHolidayc   -1.072e+03  2.563e+02   -4.183 2.87e-05 ***
## SchoolHoliday1  2.838e+02  6.729e+00   42.168 < 2e-16 ***
## StoreTypeb      2.941e+03  2.634e+01  111.652 < 2e-16 ***
## StoreTypec     -4.080e+02  7.244e+00  -56.319 < 2e-16 ***
## StoreTyped      8.570e+02  5.711e+00  150.076 < 2e-16 ***
## Assortmentb     -3.304e+03  3.465e+01  -95.373 < 2e-16 ***
## Assortmentc     5.421e+02  4.960e+00  109.291 < 2e-16 ***
## CompetitionDistance 2.862e-03  3.155e-04    9.072 < 2e-16 ***
## CompetitionOpenYears 2.300e+00  4.782e-01    4.809 1.51e-06 ***
## PopulationLevel2  1.154e+03  6.745e+00  171.034 < 2e-16 ***
## PopulationLevel3  2.427e+03  6.898e+00  351.867 < 2e-16 ***
## PopulationLevel4  4.898e+03  7.376e+00  664.044 < 2e-16 ***
## Promo1          2.304e+03  5.199e+00  443.125 < 2e-16 ***
## Promo21         -3.488e+02  1.076e+01  -32.416 < 2e-16 ***
## Promo2OpenDays   2.821e-01  5.718e-03   49.339 < 2e-16 ***
## PromoIntervalJan, Apr, Jul, Oct 2.420e+02  8.164e+00   29.649 < 2e-16 ***
## PromoIntervalMar, Jun, Sept, Dec -1.251e+00  1.040e+01   -0.120    0.904
## PromoIntervalNo Interval      NA      NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2153 on 844301 degrees of freedom
## Multiple R-squared:  0.519, Adjusted R-squared:  0.519
```

```
## F-statistic: 2.53e+04 on 36 and 844301 DF,  p-value: < 2.2e-16
```

Check some benchmarks R-square:0.519 meaning only about 52% variance of Sales can be explained by predictors included in the model, P-value shows that the linear model is meaningful anyway, but it seems rf yields a better result, hence prediction on test set will use rf model.

## 2.3 Prediction

Once we got our model, we are almost there! Just a very quick thing left—prediction.

```
# Predict using the test set
test$prediction <- predict(rf.model2, test[, 3:15])

# Save the submission with Id and predicted Sales(Only for Open=1)
submission<- test[,c(1,16)]

#view the format
head(submission)
```

```
##      Id prediction
## 844339 1    5064.980
## 844340 2    7861.872
## 844341 3   11022.824
## 844342 4    5772.017
## 844343 5    6150.727
## 844344 6    5844.600
```

```
#Write solution to CSV file
write.csv(submission, file = 'submission.csv', row.names = F)
```

Thank you for taking your time to read this analysis, any recommendation or cureness is so welcomed!